

CPSC 471: Computer Communications

Switching and Bridging

Figures from [Computer Networks: A Systems Approach](#), version 6.02dev
(Larry L. Peterson and Bruce S. Davie)

You may not distribute/post these lecture slides without written permission
from Dr. Mike Turi, ECE Dept., California State University, Fullerton

Switches

- ⦿ Devices that interconnect links of the same type to form a larger network
- ⦿ Transfers packets from an input to one or more outputs
- ⦿ Switch has fixed number of ports (I/O)
 - Limits the number of hosts that can be connected to a switch




Figure 56

Packet Switching/Forwarding

- ⦿ Large networks can be made by interconnecting switches
 - Switched networks are more scalable
- ⦿ Switches use packet switching/forwarding
 - Receives incoming packets on one link
 - Forwards them on a different link
 - Main function of the network layer
 - How does the switch decide which output link to place a packet on?
- ⦿ Three approaches to packet forwarding

Datagram/Connectionless Approach


- Include enough info in every packet for any switch to decide how to get it to destination
 - Every packet contains the complete destination address
- Switch looks at a forwarding table
 - Easy to figure out for a small, simple, static network
 - Tough for a large, complex, dynamic network
 - Could contain multiple paths between a pair of hosts
 - This is routing

Forwarding Table

Figure 57

Forwarding Table for Switch 2
(Table 5)

| Destination | Port |
|-------------|------|
| A | 3 |
| B | 0 |
| C | 3 |
| D | 3 |
| E | 2 |
| F | 1 |
| G | 0 |
| H | 0 |



Datagram Approach continued

- ⦿ A host can send a packet anywhere, anytime
 - Any packet that arrives at a switch can be immediately forwarded
- ⦿ Sender does not know if network is capable to deliver it
 - Or if host is up and running
- ⦿ Each packet is forwarded independently
 - Packets may use different paths between source and destination
 - This is robust to switch or link failures

Virtual Circuit (VC)/Connection-Oriented Approach

- ⦿ Requires setting up a virtual connection from source to destination sending data
 - Connection setup, then data transfer
- ⦿ Must establish a “connection state” in each switch between source and destination hosts
 - Consists of entry in “VC table”

VC Table

- ⦿ Virtual Circuit Identifier (VCI)
 - Uniquely identifies connection at this switch
 - Carried in the header of packets of this connection
- ⦿ Incoming interface that VC packets arrive at switch
- ⦿ Outgoing interface that VC packets depart the switch
- ⦿ A potentially different VCI for outgoing packets

VC Approach continued

- ⦿ If a packet arrives on incoming interface with VCI value in header
 - Send it out on outgoing interface with outgoing VCI value placed in the header
- ⦿ The VCI of received packets and the interface number uniquely identify the VC
- ⦿ For a new connection, each link must assign a new VCI for the connection

Establish Connection State: PVC

- ⦿ Have a network admin configure the state
 - Permanent VC or PVC
 - The VC is permanent until the network admin deletes it
- ⦿ Burdensome for large networks with many switches

VC Network

VC Table entries for Switches 1-3 (Tables 6-8)

| Switch # | Incoming Interface | Incoming VCI | Outgoing Interface | Outgoing VCI |
|----------|--------------------|--------------|--------------------|--------------|
| 1 | 2 | 5 | 1 | 11 |
| 2 | 3 | 11 | 2 | 7 |
| 3 | 0 | 7 | 1 | 4 |




Figure 59

Establish Connection State: SVC

- ⦿ A host can send messages into the network to cause the state to be established
 - Switched VC or SVC
 - This is “signaling” (resultant VC is “switched”)
 - A host may set up and delete a VC dynamically
 - Without a network admin

Signaling Example




Figure 59

Summary of Signaling

- ⦿ Sender sends a setup message
- ⦿ Each switch receives this and sets
 - Incoming and outgoing interfaces
 - Incoming VCI (picks this value)
 - Outgoing VCI is picked by next switch down the line
- ⦿ Receiver sends an ACK
- ⦿ Via ACK, each switch sets the outgoing VCI
- ⦿ Sender receives ACK

Teardown

- ⦿ When sender no longer wishes to send data to receiver
 - Sends a teardown message
- ⦿ Each switch that receives this message
 - Forwards it
 - Deletes the VC entry from its VC table

VC vs. Datagram

- ⦿ Per-packet overhead of VC is less than datagram
 - Less addressing in header
- ⦿ When the host receives a connection-setup ACK
 - It knows receiver is ready
 - It knows there is a route to the receiver
- ⦿ If switch or link fails, connection is broken
 - A new connection must be set up
 - Old connection must be torn down
 - To clean up VC tables

VC Example: X.25

- ⦿ Buffers allocated to each VC when initialized
- ⦿ Sliding window protocol run between nodes
- ⦿ VC is rejected if node does not have buffer space when connection request received
- ⦿ These are usually called hop-by-hop flow control
 - No connection establishment for datagram
 - Each switch processes each packet independently
 - Each arriving packet competes for buffer space against other packets
- ⦿ Congestion vs. Contention?

Congestion vs. Contention

⦿ Contention

- Packets queued at switch since they are competing for same output link

⦿ Congestion

- Switch has more packets than buffer space and must drop packets

⦿ X.25 does not encounter congestion

- Efficient use of buffer space?

⦿ Datagram invites congestion

VC Example: ATM

- ⦿ Asynchronous Transfer Mode (ATM)
 - Generic Flow Control (GFC)
 - Virtual Path Identifier (VPI)
 - Virtual Circuit Identifier (VCI)
 - Header Error Check (HEC) [8-bit CRC]
- ⦿ Only one size: 53 bytes (48 bytes of data)
 - Makes switch hardware simple and efficient
 - What cell size should be picked?
- ⦿ Why did ATM fail?




Figure 61

Source Routing Approach

- Info about network topology is provided by the source host
- Use an ordered list of switch ports to route packet


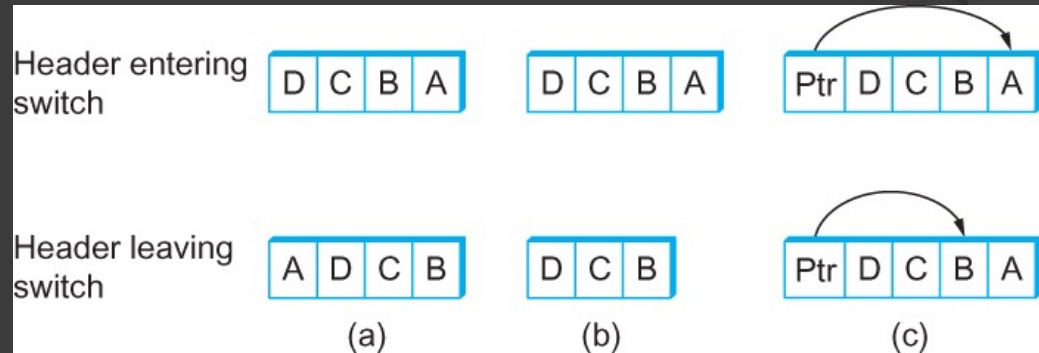


Figure 62

Source Routing continued

Figure 63



- Switch could alter header:
- Assumes the switch knows the network topology
 - Similar to building forwarding tables or figuring out where to send a setup packet
- Cannot predict how big the header will be

Strict and Loose Source Routing

⦿ Strict

- Every node along the path must be specified

⦿ Loose

- Only specifies a set of nodes to be traversed
 - Doesn't specify exactly how to get from one node to the next
- Can be helpful to limit the amount of info the source node must gather
 - Hard to get the complete path info in a large network

Bridges and LAN Switches

- ⦿ Forward packets between LANs
- ⦿ Could use a repeater between two different Ethernet segments
 - But what if this exceeds the physical limitations?
- ⦿ Put a node with a pair of Ethernet adaptors between the two Ethernets
 - This is a bridge (or a LAN switch)

Ethernet example continued

- ⦿ Bridge forwards frames from one Ethernet to another
 - Repeater operates on bits and blindly copies bits
- ⦿ Bridge implements Ethernet's collision detection and media access protocols on each interface
- ⦿ Accepts all frames and forwards them to the other Ethernet

Bridges, more generally

- Collection of LANs connected by one or more bridges is an extended LAN
- Originally bridge accepts LAN frames on inputs and forwards them to all other outputs
- A bridge is a switch
- Provides a way to increase total bandwidth of the network


Learning Bridges

- A bridge doesn't need to forward all frames it receives
- Bridge can learn and build a forwarding table

Figure 64

Forwarding Table for Bridge
(Table 9)

| Host | Port |
|------|------|
| A | 1 |
| B | 1 |
| C | 1 |
| X | 2 |
| Y | 2 |
| Z | 2 |



Learning Bridges continued

- ⦿ When a bridge first boots, the table is empty
 - Entries are added over time as frames are transmitted
 - A timeout is associated with each entry, why?
 - If a bridge receives a frame for a host not in the table, it forwards the frame out on all other ports
- ⦿ Table is an optimization, not required for correctness

Learning Bridge Example

- Consider the following topology:



- If all forwarding tables are initially empty, determine the forwarding tables for bridges B1-B4 after:
 - D sends to C
 - C sends to D
 - A sends to C

Solution to Learning Bridge Example

- ⦿ All bridges see the packet from D to C
- ⦿ Only B2-B4 see the packet from C to D
- ⦿ Only B1-B3 see the packet from A to C

- ⦿ B1: A-interface: A B2-interface: D (not C)
- ⦿ B2: B1-interface: A B3-interface: C
 B4-interface: D
- ⦿ B3: C-interface: C B2-interface: A, D
- ⦿ B4: D-interface: D B2-interface: C (not A)

Spanning Tree Algorithm

- Learning bridges will fail if the extended LAN has a loop
 - How?
- Why would an extended LAN have a loop?
- Need a different algorithm
 - Spanning Tree




Figure 65

Spanning Tree

- A subgraph of the extended LAN's graph which covers/spans all vertices but contains no loops/cycles
- Throws out some edges of the graph




Figure 66. Cyclic Graph and Spanning Tree

Spanning Tree Algorithm

- ⦿ Protocol used by bridges to agree upon a spanning tree for an extended LAN
 - Each bridge decides which ports it is and isn't willing to forward frames
 - A bridge may not participate in forwarding frames
 - Algorithm is dynamic and bridges may reconfigure themselves if a bridge fails

Spanning Tree, step-by-step

- ⦿ Each bridge has a unique ID
- ⦿ Bridge with smallest ID is the root bridge
 - Root bridge always forwards frames out over all its ports
- ⦿ Each bridge computes shortest path to root
 - This is bridge's preferred path to root
- ⦿ All bridges connected to a common LAN elect a single designated bridge that will forward frames to the root

Spanning Tree, step-by-step continued

- ⦿ Electing a designated bridge
 - Designates the bridge closest to the root
 - Use smallest ID to break ties
 - Each bridge participates in election for each LAN it is connected to
- ⦿ The bridge forwards frames over those ports which it is the dedicated bridge

Spanning Tree Algorithm Example




Figure 67

Spanning Tree Configuration

- ⦿ Bridges must exchange configuration messages
 - Decide if they are the root or a designated bridge
- ⦿ Configuration messages contain:
 - ID of bridge sending the message
 - ID of what the sending bridge believes to be the root bridge
 - The distance, in hops, from the sending bridge to the root bridge

Spanning Tree Configuration continued

- ⦿ Initially each bridge thinks it's the root
 - Sends message identifying itself as root and distance to root = 0
- ⦿ Each bridge records the current “best” configuration message it has sent or received
 - Better configuration message is:
 - Identifies a root with smaller ID
 - Identifies a root with equal ID, but shorter distance
 - Root ID and distance are equal, but sending bridge has a smaller ID

Spanning Tree Configuration continued

- ⦿ If new configuration message better than current
 - Bridge discards current message and saves the new message
 - Adds 1 to the distance-to-root field
- ⦿ Bridge stops generating configuration messages (but still forwards) if it's not the root
- ⦿ Bridge stops forwarding on a port if it's not the designated bridge for that port

Spanning Tree Configuration continued

- ⦿ When system stabilizes, only root bridge generates configuration messages
 - Designated bridges forward these
- ⦿ Spanning tree has been built, and all bridges agree which ports will be used for the spanning tree
 - Only these ports forward packets in the extended LAN

Spanning Tree Algorithm Example




Figure 67

Spanning Tree Algorithm continued

- ⦿ Root bridge continues to send configuration messages periodically
 - If a bridge fails, downstream bridges do not get these messages
 - A timeout occurs and the algorithm restarts
- ⦿ Algorithm can reconfigure the spanning tree when a bridge fails
 - Cannot forward frames over alternate paths to route around a congested bridge
 - Why is this?

Broadcast and Multicast

- ⦿ Bridge forwards a broadcast message to all other ports
- ⦿ Multicast can do this too
 - Each host decides whether to receive the message
 - Could be done better
 - What if no multicast recipients are on the other LAN?

Limited Scalability of Bridges

- ⦿ Not very scalable
 - Spanning Tree algorithm scales linearly
 - Cannot impose a hierarchy on extended LAN
 - Bridges forward all broadcast frames
 - Not every host on extended LAN may want the message
 - Broadcast does not scale
- ⦿ Increase scalability by using virtual LANs

Virtual LANs (VLANs)

- Partition a single extended LAN to be into several (seemingly separate) LANs
- Each LAN is assigned an identifier
- Packets travel from one segment to another if they have the same identifier




Figure 68

VLANs continued

- ⦿ Limits the number of segments that receive a broadcast message
- ⦿ Need to configure a VLAN ID on each port of the bridges
- ⦿ Bridge observes which port a packet arrives on and which VLAN that port is configured for
- ⦿ Possible to change logical topology without moving wires or changing addresses

Bridges Limited in Network Interconnection

- ⦿ Meant to connect similar LANs
 - Use the network's frame header
 - Support networks with exactly the same format for addresses
 - E.g., Ethernets to Ethernets, etc.
 - Bridge between Ethernet and 802.11 is okay
 - Both use same 48-bit address format
- ⦿ Do not generalize to other kinds of networks with different addressing formats
 - E.g., ATM

Bridges allow Transparent Connection of Multiple LANs

- ⦿ Networks can be connected without end hosts having to run (or be aware of) any additional protocols
- ⦿ Not safe to design network software assuming it will run on a single LAN
 - If a bridge becomes congestion, it'll drop frames
 - Ethernet rarely drops frames
 - Latency between a pair of hosts on an extended LAN can become larger and more variable
 - Ethernet's latency is small and predictable
 - Frames may be reordered in an extended LAN
 - Frames are never ordered on a single Ethernet